

EECS 489

Computer Networks

Winter 2025

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Recap: End-to-end principle

- ❑ Functions that can be *completely* and *correctly* implemented *only* with the knowledge of application end host, should not be pushed into the network
- ❑ **Fate sharing**: fail together or don't fail at all
 - “it is acceptable to lose the state information associated with an entity if, at the same time, the entity itself is lost”

Agenda

- HTTP and the Web
- Improving HTTP Performance

The Web: Precursor

- ❑ 1945, [Vannevar Bush](#), Memex
 - Concept of the web based on microfilms
- ❑ 1967, [Ted Nelson](#), Project Xanadu
 - A world-wide publishing network to store information as connected literature
 - Coined the term “Hypertext”
- ❑ 1968, [Douglas Engelbart](#), NLS (oN-Line System)
 - The mother of all demos

The Web: History

- World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)
 - First HTTP implementation – 1990
 - » **Tim Berners-Lee** at CERN
 - HTTP/0.9 – 1991
 - » Simple GET command for the Web
 - HTTP/1.0 – 1992
 - » Client/server information, simple caching



AWARD WINNER

Sir Tim Berners-Lee

ACM A. M. Turing Award (2016)

ACM Software System Award (1995)

2016 ACM A.M. Turing Award

The Web: History (cont'd)

- World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)
 - HTTP/1.1 – 1996
 - » Performance and security optimizations
 - HTTP/2 – 2015
 - » Latency optimizations via request multiplexing over single TCP connection
 - » Binary protocol instead of text
 - » Server push

The Web: History (cont'd)

- World Wide Web (WWW): a distributed database of “pages” linked through Hypertext Transport Protocol (HTTP)
 - HTTP/3 – 2022
 - » Built on top of QUIC, which is a user-space congestion control protocol on top of UDP
 - » Solves the **head-of-line (HOL) blocking problem** when multiplexing over a single TCP connection

What does it consist of?

- ❑ Who uses it?
- ❑ Who provides the content?
- ❑ How do they communicate?

- ❑ How do we find the content?
- ❑ How is the content organized?
- ❑ How is it displayed?

Web components

❑ Infrastructure:


- Clients
- Servers (DNS, CDN, Datacenters)

❑ Content:


- URL: naming content
- HTML: formatting content

❑ Protocol for exchanging information: HTTP


Why is there nothing about the network?



What we want



What we get



URL: Uniform Record Locator

- ❑ `protocol://host-name[:port]/directory-path/resource`
- ❑ Extend the idea of hierarchical hostnames to include anything in a file system
 - `https://github.com/mosharaf/eecs489/blob/f24/slides/012224.pptx`
- ❑ Extend to program executions as well...
 - `https://www.google.com/search?q=eecs489`
 - Server-side processing can be included in the name


URL: Uniform Record Locator

- `protocol://host-name[:port]/directory-path/resource`
 - `protocol`: http, ftp, https, smtp, rtsp, *etc.*
 - `host-name`: DNS name, IP address
 - `port`: defaults to protocol's standard port
 - » *E.g.*, http: 80, https: 443
 - `directory path`: hierarchical, reflecting file system
 - `resource`: Identifies the desired resource

Hyper Text Transfer Protocol (HTTP)

- ❑ Client-server architecture
 - Server is “always on” and “well known”
 - Clients initiate contact to server
- ❑ Synchronous request/reply protocol
 - Runs over TCP, Port 80
- ❑ **Stateless**
- ❑ ASCII format before HTTP/2 but Binary since then

Steps in canonical HTTP request/response



Method types (HTTP 1.1)

- GET, HEAD

- POST

 - Send information (e.g., web forms)

- PUT

 - Uploads file in entity body to path specified in URL field

- DELETE

 - Deletes file specified in the URL field

Client-to-server communication

? HTTP Request Message

- Request line: method, resource, and protocol version

request line → GET /somedir/page.html HTTP/1.1


header lines → Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
(blank line)

*carriage return line feed
indicates end of message*

Client-to-server communication

HTTP Request Message

- Request line: method, resource, and protocol version
- Request headers: provide info or modify request
- Body: optional data (e.g., to “POST” data to server)



Server-to-client communication

? HTTP Response Message

- Status line: protocol version, status code, status phrase
- Response headers: provide information
- Body: optional data

status line

(protocol, status code, status phrase)

header lines

data

e.g., requested HTML file

HTTP/1.1 200 OK

Connection close

Date: Thu, 06 Jan 2017 12:00:15 GMT

Server: Apache/1.3.0 (Unix)

Last-Modified: Mon, 22 Jun 2006 ...

Content-Length: 6821

Content-Type: text/html

(blank line)

data data data data data ...


HTTP is stateless

- Each request-response treated **independently**
 - Servers not required to retain state
- **Good:** Improves scalability on the server-side
 - Failure handling is easier
 - Can handle higher rate of requests
 - Order of requests doesn't matter
- **Bad:** Some applications need persistent state
 - Need to uniquely identify user or store temporary info
 - e.g., Shopping cart, user profiles, usage tracking, ...

How does a stateless protocol keep state?

State in a stateless protocol: Cookies

- **Client-side** state maintenance
 - Client stores small state on behalf of server
 - Client sends state in future requests to the server
- Can provide authentication



Beyond cookies

- ❑ Cookies provide excellent marketing opportunities and create concerns for privacy
 - Advertising companies tracks your preferences and viewing history across sites
- ❑ Many are trying to curtail the (mis)use of cookies
 - Example: Google Chrome was planning to deprecate third-party cookies

5-MINUTE BREAK!

Announcements

- Get started on Assignment 1

Performance goals

☐ User

- Fast downloads (not identical to low-latency communication!)
- High availability

☐ Content provider

- Happy users (hence, above)
- Cost-effective infrastructure

☐ Network (secondary)

- Avoid overload

Solutions?

Improve networking protocols including HTTP, TCP, etc.



? User

- Fast downloads (not identical to low-latency communication!)
- High availability

? Content provider

- Happy users (hence, above)
- Cost-effective infrastructure

? Network (secondary)

- Avoid overload

Solutions?

Improve networking protocols including HTTP, TCP, etc.

? User

- Fast downloads (not identical to low-latency communication!)
- High availability

? Content provider

- Happy users (hence, above)
- Cost-effective infrastructure

? Network (secondary)

- Avoid overload

Caching and replication

Solutions?

Improve networking protocols including HTTP, TCP, etc.

? User

- Fast downloads (not identical to low-latency communication!)
- High availability

? Content provider

- Happy users (hence, above)
- Cost-effective infrastructure

? Network (secondary)

- Avoid overload

Caching and replication

Exploit economies of scale; e.g., webhosting, CDNs, datacenters

HTTP performance

- Most Web pages have multiple objects
 - e.g., HTML file and a bunch of embedded images
- How do you retrieve those objects (naively)?
 - One item at a time
- New TCP connection per (small) object!


Object request response time

RTT (round-trip time)

- Time for a small packet to travel from client to server and back

Response time

- 1 RTT for TCP setup
- 1 RTT for HTTP request and first few bytes
- Transmission time
- **Total** = 2RTT + Transmission Time




Non-persistent connections

- ❑ Default in HTTP/1.0
- ❑ $2RTT + \Delta$ for each object in the HTML file!
 - One more $2RTT + \Delta$ for the HTML file itself
- ❑ Doing the same thing over and over again
 - Inefficient

Concurrent requests and responses

- Use multiple connections in parallel
- Does not necessarily maintain order of responses

- Client = 😊
- Content provider = 😊
- Network = 😞 Why?




Persistent connections

- Maintain TCP connection across multiple requests
 - Including transfers subsequent to current page
 - Client or server can tear down connection
- Advantages
 - Avoid overhead of connection set-up and tear-down
 - Allow underlying layers (e.g., TCP) to learn about RTT and bandwidth characteristics
- Default in HTTP/1.1

Pipelined requests & responses

- ❑ Batch requests and responses to reduce the number of packets
 - Multiple requests can be contained in one TCP segment
- ❑ Data are sent in a FIFO manner
 - Can lead to **head-of-line (HOL) blocking** if many small responses follow a large one
 - Not supported by default by major browsers circa 2015
- ❑ **Solution**
 - Priority and preemption



Scorecard: Getting n small objects

- Time dominated by latency
- One-at-a-time: $\sim 2n$ RTT
- m concurrent: $\sim 2\lceil n/m \rceil$ RTT
- Persistent: $\sim (n+1)$ RTT
- Pipelined: ~ 2 RTT
- Pipelined and Persistent: ~ 2 RTT first time; RTT later for another n from the same site

Scorecard: Getting n large objects each of size F

- Time dominated by TCP throughput B_C ($\leq B_L$), where bottleneck link bandwidth is B_L
 - Assuming all TCP connections go through the same bottleneck link
- One-at-a-time: $\sim nF/B_C$
- m concurrent: $\sim nF/(mB'_C)$
 - Assuming each TCP connection gets the same throughput (B'_C), where $mB'_C \leq B_L$
- Pipelined and/or persistent: $\sim nF/B_C$
 - The only thing that helps is higher throughput

Caching

- Why does caching work?
 - Exploits locality of reference
- How well does caching work?
 - Very well, up to a limit
 - Large overlap in content
 - But many unique requests
 - » A universal story!
 - » Effectiveness of caching grows logarithmically with size

Caching: How

- ❑ Modifier to GET requests:
 - **If-modified-since** – returns “not modified” if resource not modified since specified time

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
If-modified-since: Wed, 18 Jan 2017 10:25:50 GMT
(blank line)
```

Caching: How

- ❑ Modifier to GET requests:
 - **If-modified-since** – returns “not modified” if resource not modified since specified time
- ❑ Client specifies “**if-modified-since**” time in request
- ❑ Server compares this against “last modified” time of resource
- ❑ Server returns “Not Modified” if resource has not changed
- ❑ or a “OK” with the latest version otherwise

Caching: How

❏ Modifier to GET requests:

- **If-modified-since** – returns “not modified” if resource not modified since specified time

❏ Response header:

- **Expires** – how long it's safe to cache the resource
- **No-cache** – ignore all caches; always get resource directly from server


Caching: Where?

□ Options

- Client (browser)
- Forward proxies
- Reverse proxies
- Content Distribution Network


Caching: Where?

- Many clients transfer same information
 - Generate unnecessary server and network load
 - Clients experience unnecessary latency




Caching with Reverse Proxies

- Cache documents close to server
 - Decrease server load
 - By content provider



Caching with Forward Proxies

- Cache documents close to clients
 - Reduce network traffic and decrease latency
 - By ISPs or enterprises



Summary

□ HTTP/1.1

- Text-based protocol
- Replaced by binary HTTP/2 protocol, which has been replaced by HTTP/3 in 2022

□ Many ways to improve performance

- Pipelining and batching
- Caching in proxies and CDNs
- Datacenters