

EECS 489

Computer Networks

Winter 2025

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Agenda

- TCP congestion control wrap-up
- TCP throughput equation
- Problems with congestion control

Recap

□ Flow Control

- Restrict window to RWND to make sure that the receiver isn't overwhelmed

□ Congestion Control

- Restrict window to CWND to make sure that the network isn't overwhelmed

□ Together

- Restrict window to $\min\{\text{RWND}, \text{CWND}\}$ to make sure that neither the receiver nor the network are overwhelmed

CC Implementation

- ▣ States at sender
 - **CWND** (initialized to a small constant)
 - **ssthresh** (initialized to a large constant)
 - **Timer**
- ▣ Events
 - **ACK** (new data)
 - **Timeout**

Event: ACK (new data)

□ If $CWND < ssthresh$

➤ $CWND += 1$

- *$CWND$ packets per RTT*
- *Hence, after one RTT with no drops:
 $CWND = 2 \times CWND$*

Event: ACK (new data)

□ If $CWND < ssthresh$

➤ $CWND += 1$

Slow start phase

□ Else

➤ $CWND = CWND + 1/CWND$


Congestion avoidance phase

- *CWND packets per RTT*
- *Hence, after one RTT with no drops:*
 $CWND = CWND + 1$

Event: TimeOut

- On Timeout
 - $\text{ssthresh} \leftarrow \text{CWND}/2$
 - $\text{CWND} \leftarrow 1$

AIMD leads to TCP sawtooth



CC Implementation (cont.)

▣ States at sender

- CWND (initialized to a small constant)
- ssthresh (initialized to a large constant)
- Timer
- dupACKcount

▣ Events


- ACK (new data)
- Timeout
- dupACK (duplicate ACK for old data)

Event: dupACK

- ▣ dupACKcount ++
- ▣ If dupACKcount = 3 /* fast retransmit */
 - ssthresh = CWND/2
 - CWND = CWND/2

Example

Window




Slow-start restart: Go back to $\text{CWND} = 1 \text{ MSS}$, but take advantage of knowing the previous value of CWND

Why AIMD?

- Recall the three issues
 - Finding available bottleneck bandwidth
 - Adjusting to bandwidth variations
 - Sharing bandwidth
- Two goals for bandwidth sharing
 - **Efficiency**: High utilization of link bandwidth
 - **Fairness**: Each flow gets equal share

AIMD is fair and efficient

- ? Increase: $x + a_i$
- ? Decrease: $x * b_D$
- ? Converges to fairness



Not done yet!

- **Problem:** congestion avoidance too slow in recovering from an isolated loss

Example

- Consider a TCP connection with:
 - CWND=10 packets
 - Last ACK was for packet # 101
 - »i.e., receiver expecting next packet to have seq. no. 101
- 10 packets [101, 102, 103,..., 110] are in flight
 - Packet 101 is dropped

Timeline: [~~101~~, 102, ..., 110]

- ACK 101 (due to 102) $cwnd=10$ dupACK#1 (no xmit)
- ACK 101 (due to 103) $cwnd=10$ dupACK#2 (no xmit)
- ACK 101 (due to 104) $cwnd=10$ dupACK#3 (no xmit)
- RETRANSMIT 101 $ssthresh=5$ $cwnd=5$
- ACK 101 (due to 105) $cwnd=5 + 1/5$ (no xmit)
- ACK 101 (due to 106) $cwnd=5 + 2/5$ (no xmit)
- ACK 101 (due to 107) $cwnd=5 + 3/5$ (no xmit)
- ACK 101 (due to 108) $cwnd=5 + 4/5$ (no xmit)
- ACK 101 (due to 109) $cwnd=5 + 5/5$ (no xmit)
- ACK 101 (due to 110) $cwnd=6 + 1/6$ (no xmit)
- ACK 111 (due to 101) ← only now can we transmit new packets
- Plus no packets in flight so ACK “clocking” (to increase CWND) stalls for another RTT

Solution: Fast recovery

- Idea: Grant the sender temporary “credit” for each dupACK so as to keep packets in flight
- If dupACKcount = 3
 - ssthresh = CWND/2
 - CWND = ssthresh + 3
- While in fast recovery
 - CWND = CWND + 1 for each additional dupACK
- Exit fast recovery after receiving new ACK
 - set CWND = ssthresh


Example

- Consider a TCP connection with:
 - CWND=10 packets
 - Last ACK was for packet # 101
 - »i.e., receiver expecting next packet to have seq. no. 101
- 10 packets [101, 102, 103,..., 110] are in flight
 - Packet 101 is dropped


Timeline: [~~101~~, 102, ..., 110]

- ACK 101 (due to 102) cwnd=10 dup#1
- ACK 101 (due to 103) cwnd=10 dup#2
- ACK 101 (due to 104) cwnd=10 dup#3
- RETRANSMIT 101 ssthresh=5 cwnd= 8 (5+3)
- ACK 101 (due to 105) cwnd= 9 (no xmit)
- ACK 101 (due to 106) cwnd=10 (no xmit)
- ACK 101 (due to 107) cwnd=11 (xmit 111)
- ACK 101 (due to 108) cwnd=12 (xmit 112)
- ACK 101 (due to 109) cwnd=13 (xmit 113)
- ACK 101 (due to 110) cwnd=14 (xmit 114)
- ACK 111 (due to 101) cwnd = 5 (xmit 115) ← exiting fast recovery
- Packets 111-114 already in flight
- ACK 112 (due to 111) cwnd = $5 + 1/5$ ← back in cong. avoidance


TCP state machine




Timeouts → Slow Start




dupACKs → Fast Recovery



New ACK changes state ONLY from Fast Recovery



TCP state machine



TCP flavors

❑ TCP-Tahoe

- CWND = 1 on 3 dupACKs

❑ TCP-Reno

- CWND = 1 on timeout
- CWND = CWND/2 on 3 dupACKs

❑ TCP-newReno

- TCP-Reno + improved fast recovery

❑ TCP-SACK

- Incorporates selective acknowledgements

**Our default
assumption**


How can they coexist?

- All follow the same principle
 - Increase CWND on good news
 - Decrease CWND on bad news


5-MINUTE BREAK!

TCP THROUGHPUT EQUATION

A simple model for TCP throughput




A simple model for TCP throughput



Implications (1): Different RTTs

$$\text{Throughput} = \sqrt{\frac{3}{2}} \frac{1}{RTT \sqrt{p}}$$

- Flows get throughput inversely proportional to RTT
- TCP unfair in the face of heterogeneous RTTs!



Implications (2): High-speed TCP

$$\text{Throughput} = \sqrt{\frac{3}{2}} \frac{1}{RTT \sqrt{p}}$$

- Assume $RTT = 100\text{ms}$, $MSS=1500\text{bytes}$,
 $BW=100\text{Gbps}$
- What value of p is required to reach 100Gbps throughput?
 - $\sim 2 \times 10^{-12}$
- How long between drops?
 - ~ 16.6 hours
- How much data has been sent in this time?
 - ~ 6 petabits

Implications (3): Rate-based CC

$$\text{Throughput} = \sqrt{\frac{3}{2}} \frac{1}{RTT \sqrt{p}}$$

- ❑ TCP throughput swings between $W/2$ to W
- ❑ Apps may prefer steady rates (e.g., streaming)
- ❑ “Equation-Based Congestion Control”
 - Ignore TCP’s increase/decrease rules and just follow the equation
 - Measure drop percentage p , and set rate accordingly
- ❑ Following the TCP equation ensures “TCP friendliness”
 - i.e., use no more than TCP does in similar setting

Implications (4): Loss not due to congestion?

- ❑ TCP will confuse corruption with congestion
- ❑ Flow will cut its rate
 - Throughput $\sim 1/\sqrt{p}$ where p is loss prob.
 - Applies even for non-congestion losses!

Implications (5):

Short flows cannot ramp up

- 50% of flows have $< 1500\text{B}$ to send; 80% $< 100\text{KB}$
- Implications
 - Short flows never leave slow start!
 - » They never attain their fair share
 - Too few packets to trigger dupACKs
 - » Isolated loss may lead to timeouts
 - » At typical timeout values of $\sim 500\text{ms}$, might severely impact flow completion time

Implications (6):

Short flows share long delays

- A flow deliberately overshoots capacity, until it experiences a drop
- Means that delays are large, and are large for everyone
 - Consider a flow transferring a 10GB file sharing a bottleneck link with 10 flows transferring 100B
 - Larger flows dominate smaller ones

Implications (7): Cheating

- Three easy ways to cheat
 - Increasing CWND faster than $+1 \text{ MSS per RTT}$

Implications (7): Cheating

- Three easy ways to cheat
 - Increasing CWND faster than +1 MSS per RTT
 - Using large initial CWND
 - » Common practice by many companies

Implications (7): Cheating

- Three easy ways to cheat
 - Increasing CWND faster than $+1 \text{ MSS per RTT}$
 - Using large initial CWND
 - » Common practice by many companies
 - Opening many connections

Open many connections



| Assume

- A starts 10 connections to B
- D starts 1 connection to E
- Each connection gets about the same throughput

| Then A gets 10 times more throughput than D

Implications (8): CC intertwined with reliability

- | CWND adjusted based on ACKs and timeouts
- | Cumulative ACKs and fast retransmit/recovery rules
- | Complicates evolution
 - Changing from cumulative to selective ACKs is hard
- | Sometimes we want CC but not reliability
 - e.g., real-time applications
- | We may also want reliability without CC

Recap: TCP problems

Routers tell endpoints
if they're congested

- | Misled by non-congestion losses
- | Fills up queues leading to high delays
- | Short flows complete before discovering available capacity
- | AIMD impractical for high speed links
- | Saw tooth discovery too choppy for some apps
- | Unfair under heterogeneous RTTs
- | Tight coupling with reliability mechanisms
- | End hosts can cheat

Routers tell
endpoints what
rate to send at

Routers enforce
fair sharing

Could fix many of these with some help from routers!

Summary

- | TCP works even though it has many flaws
- | Many of them can be fixed via assistance from the network
 - Starting from next lecture: The Network Layer
- | NO lecture on Feb 12!



Adapting TCP to high speed

- | Once past a threshold speed, increase CWND faster
 - A proposed standard [Floyd'03]: once speed is past some threshold, change equation to $p^{-.8}$ rather than $p^{-.5}$
 - Let the additive constant in AIMD depend on CWND
- | Other approaches?
 - Multiple simultaneous connections ([hack but works today](#))
 - Router-assisted approaches