

EECS 489

Computer Networks

Winter 2025

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Logistics

- You're ONLY allowed one double-sided 8.5in x 11in note sheet
 - Basic calculators are allowed but you won't need it

General guidelines (1)

- ❑ Test only assumes material covered in lecture, discussion sections, quizzes, and assignments
 - Text: only to clarify details and context for the above
- ❑ The test doesn't require you to do complicated calculations
 - Use this as a hint to determine if you're on right track
- ❑ You don't need to memorize anything
- ❑ You do need to understand how things work

General guidelines (2)

□ Be prepared to:

- Weigh design options outside of the context we studied them in
- Contemplate new designs we haven't covered in detail but can be put together
 - »e.g., I introduce a new IP address format; how does this affect..”
- Reason from what you know about the pros/cons of solutions we did study

General guidelines (3)

- Exam format
 - MCQ questions
 - Networking use cases
 - » Questions not ordered in terms of complexity

General guidelines (4)

- 75-minute midterm exam
 - Friday, Feb. 21 from 7PM - 8:30PM
 - » Room assignments (by the first two letters of your username) are as follows:
 - AA - DB: Chrysler 133
 - DC - ZZ: Chrysler 220
 - SSD accommodations
 - » Friday, Feb. 21 from 6PM at DOW 2150
 - Other conflicts
 - » You have received separate emails by now

What's covered in the midterm?

- ▣ Lectures: 1—9 (Up to “More Congestion Control”)
- ▣ Assignments: 1
- ▣ Discussion sections: 1—6 (Content related to lectures 1—9 and assignment 1)

This review

- ❑ Walk through what you're expected to know at this point: key topics, important aspects of each
- ❑ Not covered in review **does NOT imply** you don't need to know it
 - But if it's covered today, you should know it
- ❑ Summarize, not explain
 - Stop me when you want to discuss something further!

Topics

- ▣ Basics (lectures 1–2)
- ▣ Application layer (lectures 3–5)
 - HTTP, DNS, CDN, Video Streaming, and Cloud
- ▣ Transport layer (lectures 6–9)
 - UDP vs. TCP
 - TCP details: reliability and flow control
 - TCP congestion control: general concepts only

Basic concepts

- You should know:
 - Packet vs. circuit switching
 - Statistical multiplexing
 - Link characteristics
 - Packet delays

Switched networks

- End-systems and networks connected by switches instead of directly connecting them
- Allows us to **scale**
 - For example, directly connecting N nodes to each other would require N^2 links!

Two approaches to sharing

□ Packet switching

- Network resources consumed on demand per-packet
- Admission control: per packet

□ Circuit switching


- Network resources reserved a priori at “connection” initiation
- Admission control: per connection

Statistical multiplexing


- Allowing more demands than the network can handle
 - Hoping that not all demands are required at the same time
 - Good for bursty traffic (average $\ll$ peak demand)
 - Packet switching exploits statistical multiplexing better than circuit switching

Delay

□ Consists of four components

- Transmission delay
 - Propagation delay
 - Queuing delay
 - Processing delay
- 
- due to link properties
- due to traffic mix and switch internals

End-to-end delay



What we want


`http://123.xyz`




123.xyz server



(Some of) What happens...



(More of) What happens



What we get



123.xyz server




Layers


- ▣ Layer: a part of a system with well-defined interfaces to other parts
- ▣ One layer interacts only with layer above and layer below
- ▣ Two layers interact only through the interface between them

Layers in practice


- Lower three layers implemented everywhere
- Top two layers implemented only at hosts



Layer encapsulation: Protocol headers



IP is the narrow waist of the layering hourglass



Topics

- ▣ Basics (lectures 1–2)
- ▣ Application layer (lectures 3–5)
 - HTTP, DNS, CDN, Video Streaming, and Cloud
- ▣ Transport layer (lectures 6–9)
 - UDP vs. TCP
 - TCP details: reliability and flow control
 - TCP congestion control: general concepts only

Hyper Text Transfer Protocol (HTTP)

- ❑ Client-server architecture
 - Server is “always on” and “well known”
 - Clients initiate contact to server
- ❑ Synchronous request/reply protocol
 - Runs over TCP, Port 80
- ❑ Stateless
- ❑ ASCII format
 - Before HTTP/2


Object request response time

RTT (round-trip time)

- Time for a small packet to travel from client to server and back

Response time

- 1 RTT for TCP setup
- 1 RTT for HTTP request and first few bytes
- Transmission time
- **Total** = 2RTT + Transmission Time



Improving HTTP performance

- Optimizing connections using three “P”s
 - Persistent connections
 - Parallel/concurrent connections
 - Pipelined transfers over the same connection
- Caching
 - Forward proxy: close to clients
 - Reverse proxy: close to servers
- Replication

Scorecard: Getting n small objects

- | Time dominated by latency
- | One-at-a-time: $\sim 2n$ RTT
- | m concurrent: $\sim 2\lceil n/m \rceil$ RTT
- | Persistent: $\sim (n+1)$ RTT
- | Pipelined: ~ 2 RTT
- | Pipelined and Persistent: ~ 2 RTT first time;
RTT later for another n from the same site

Scorecard: Getting n large objects each of size F

- | Time dominated by TCP throughput B_C ($\leq B_L$), where bottleneck link bandwidth is B_L
 - Assuming all TCP connections go through the same bottleneck link
- | One-at-a-time: $\sim nF/B_C$
- | m concurrent: $\sim nF/(mB'_C)$
 - Assuming each TCP connection gets the same throughput (B'_C), where $mB'_C \leq B_L$
- | Pipelined and/or persistent: $\sim nF/B_C$
 - The only thing that helps is higher throughput


Content Distribution Networks (CDN)

- | Caching and replication as a service
- | Combination of caching and replication
 - **Pull**: Direct result of clients' requests (caching)
 - **Push**: Expectation of high access rate (replication)


Hierarchies in the DNS

- | Three intertwined hierarchies
 - Hierarchical namespace
 - »As opposed to flat namespace
 - Hierarchically administered
 - »As opposed to centralized
 - (Distributed) hierarchy of servers
 - »As opposed to centralized storage


Name resolution




Name resolution




Name resolution




Name resolution



Name resolution: Recursive



Name resolution: Iterative



DNS caching

- | Performing all these queries takes time
 - Up to 1-second latency before starting download
- | Caching can greatly reduce overhead
 - The top-level servers very rarely change
 - Popular sites (e.g., www.google.com) visited often
 - Local DNS server often has the information cached
- | How DNS caching works
 - DNS servers cache responses to queries
 - Responses include a “time to live” (TTL) field
 - Server deletes cached entry after TTL expires

5-MINUTE BREAK!

Announcements

- | Please fill up midterm teaching evaluation

HTTP streaming

- | Video is stored at an HTTP server with a URL
- | Clients send a GET request for the URL
- | Server sends the video file as a stream
- | Client first buffers for a while to minimize interruptions later
- | Once the buffer reaches a threshold
 - The video plays in the foreground
 - More frames are downloaded in the background


DASH : Dynamic Adaptive Streaming over HTTP

- | Keep multiple resolutions of the same video
 - Stored in a manifest file in the HTTP server
- | Client asks for the manifest file first to learn about the options
- | Asks for chunks at a time and measures available bandwidth while they are downloaded
 - Low bandwidth $\Rightarrow$ switch to lower bitrate
 - High bandwidth $\Rightarrow$ switch to higher bitrate

Applications

- | Common theme: **parallelism**
 - Applications decomposed into tasks
 - Running in parallel on different machines
- | Two common paradigms
 - Partition-Aggregate
 - Map-Reduce


Traditional datacenter networks



Challenges


- | Not enough bandwidth
 - **Oversubscription**: Less bandwidth in the ToR-Agg links than all the servers' bandwidth in the rack
 - **Oversubscription ratio**: Ratio between bandwidth underneath and bandwidth above
- | Not enough paths between server pairs
 - Load balancing issues
 - Failure recovery issues

Modern datacenter networks: More bandwidth, more paths



Clos topology

- | Multi-stage network
- | k pods, where each pod has two layers of $k/2$ switches
 - $k/2$ ports up and $k/2$ down
- | All links have the same b/w
- | At most $k^3/4$ machines
- | Example
 - $k = 4$
 - 16 machines
- | For $k=48$, 27648 machines



Topics

- | Basics (lectures 1–2)
- | Application layer (lectures 3–5)
 - HTTP, DNS, CDN, Video Streaming, and Cloud
- | **Transport layer (lectures 6–9)**
 - UDP vs. TCP
 - TCP details: reliability and flow control
 - TCP congestion control: general concepts only

Role of the transport layer

- | (1) Communication between application processes
 - Mux and demux from/to application processes
 - Implemented using ports
- | (2) Provide common end-to-end services for app layer
 - Reliable, in-order data delivery
 - Well-paced data delivery

UDP vs. TCP

- Both UDP and TCP perform mux/demux via ports

	UDP	TCP
Data abstraction	Packets (datagrams)	Stream of bytes of arbitrary length
Service	Best-effort (same as IP)	•Reliability •In-order delivery •Congestion control •Flow control

Reliable transport: General concepts

- | Checksums (for error detection)
- | Timers (for loss detection)
- | Acknowledgments (feedback from receiver)
 - Cumulative: “received everything up to X”
 - Selective: “received X”
- | Sequence no (detect duplicates, accounting)
- | Sliding windows (for efficiency)

You should know:

- what these concepts are
- why they exist
- how TCP uses them

Designing a reliable transport protocol

- | **Stop and wait** is correct but inefficient
 - Works packet by packet (of size DATA)
 - Throughput is $(\text{DATA} / \text{RTT})$
- | **Sliding window**: use pipelining to increase throughput
 - n packets at a time results in higher throughput
 - $\text{MIN}(n * \text{DATA} / \text{RTT}, \text{Link Bandwidth})$

The TCP abstraction

- | TCP delivers a reliable, in-order, byte stream
- | **Reliable**: TCP resends lost packets (recursively)
 - Until it gives up and shuts down connection
- | **In-order**: TCP only hands consecutive chunks of data to application
- | **Byte stream**: TCP assumes there is an incoming stream of data, and attempts to deliver it to app

Things to know about TCP


- | How TCP achieves reliability
 - | RTT estimation
 - | Connection establishment/teardown
 - | Flow Control
 - | Congestion Control (concepts only)
-
- | For each, know how the functionality is implemented and why it is needed

Reliability

- | Having TCP take care of it simplifies application development
- | How
 - Checksums and timers (for error and loss detection)
 - Fast retransmit (to detect faster-than-timeout loss)
 - Cumulative ACKs (receiver feedback: what's lost?)
 - Sliding windows (for efficiency)
 - Buffers at sender (hold packets until ACKs arrive)
 - Buffers at receiver (to reorder packets before delivery to application)

Establishing/terminating a TCP connection

- | Three-way handshake to establish connection
 - Host A sends a SYN (open; “synchronize sequence numbers”) to host B
 - Host B returns a SYN acknowledgment (SYN ACK)
 - Host A sends an ACK to acknowledge the SYN ACK
- | Three-way handshake to terminate (normal operation)



Flow control

| Why?

- TCP at the receiver must buffer a packet until all packets before it (in byte-order) have arrived and the receiving application has consumed available bytes
- Hence, receiver advances its window when the receiving application consumes data
- Sender advances its window when new data ACK'd
- Risk of sender overrunning the receiver's buffers

| How?

- “Advertised Window” field in TCP header

Congestion control

| Why?

- Because the network itself can be the bottleneck
- Should make efficient use of available network capacity
 - » While sharing available capacity fairly with other flows
 - » And adapting to changes in available capacity

| How?

- Dynamically adapts the size of the sending window

Put together

- | Flow Control

- Restrict window to RWND to make sure that the receiver isn't overwhelmed

- | Congestion Control

- Restrict window to CWND to make sure that the network isn't overwhelmed

- | Together

- Restrict window to $\min\{\text{RWND}, \text{CWND}\}$ to make sure that neither the receiver nor the network are overwhelmed

CC implementation

- | States at sender
 - **CWND** (initialized to a small constant)
 - **ssthresh** (initialized to a large constant)
 - **dupACKcount** and **timer**
- | Events
 - **ACK** (new data)
 - **dupACK** (duplicate ACK for old data)
 - **Timeout**

Event: ACK (new data)

| If $CWND < ssthresh$

➤ $CWND += 1$

- *$CWND$ packets per RTT*
- *Hence, after one RTT with no drops:
 $CWND = 2 \times CWND$*

Event: ACK (new data)

| If $CWND < ssthresh$

➤ $CWND += 1$

Slow start phase

| Else

➤ $CWND = CWND + 1/CWND$

Congestion avoidance phase

- *CWND packets per RTT*
- *Hence, after one RTT with no drops:*
 $CWND = CWND + 1$

Event: TimeOut


- | On Timeout
 - $\text{ssthresh} \leftarrow \text{CWND}/2$
 - $\text{CWND} \leftarrow 1$

Event: dupACK

- | dupACKcount ++
- | If dupACKcount = 3 /* fast retransmit */
 - ssthresh = CWND/2
 - CWND = CWND/2

Example

Window




Slow-start restart: Go back to $\text{CWND} = 1 \text{ MSS}$, but take advantage of knowing the previous value of CWND


TCP flavors

- | TCP-Tahoe
 - CWND =1 on 3 dupACKs
- | TCP-Reno
 - CWND =1 on timeout
 - CWND = CWND/2 on 3 dupACKs
- | TCP-newReno
 - TCP-Reno + improved fast recovery
- | TCP-SACK
 - Incorporates selective acknowledgements

A simple model for TCP throughput



A simple model for TCP throughput



Summary

- | Practice Exams on GitHub